

Hallucination Mitigation in Reasoning Models

Motivation

Following the emergence of DeepSeek and GPT-5, reasoning-oriented large language models have attracted significant research attention. DeepSeek's remarkable performance stems from skipping the conventional supervised fine-tuning stage and instead employing pure GRPO-based reinforcement learning, which fundamentally enhances reasoning optimization through direct policy refinement. However, this approach has also been observed to increase the frequency of hallucinations—broadly referring to the model's tendency to generate incorrect or unsupported responses. This project aims to investigate potential mitigation strategies by refining chain-of-thought fine-tuning and reward-model design to reduce hallucination occurrence while preserving reasoning quality.

Planned Work

Overview

This project aims not only to evaluate but also to **mitigate hallucinations in reasoning-oriented LLMs** through targeted optimization of the **reward model** and **CoT fine-tuning**. Specifically, we will investigate how adjusting **PRM** and **ORM** within the **GRPO training framework** affects hallucination behavior. By dynamically balancing these components and identifying hallucination onset stages during multi-step reasoning, we aim to develop a mechanism that both detects and suppresses model “fabrication” tendencies while maintaining reasoning depth.


Methodology and Experimental Plan

The study will proceed in structured phases, combining model evaluation, dataset expansion, and iterative ablation experiments:

Phase	Description
0. Define hallucination & test scope	Establish a precise operational definition of hallucination, classify subtypes (perceptual, delusional, confabulatory, etc.), and delimit the evaluation scope across reasoning benchmarks.
1. Small-scale model + benchmark evaluation	Conduct pilot runs using a compact model to validate the dataset, prompts, and automated scoring workflow.
2. Data expansion & model improvement	Expand and refine the dataset, perform fine-tuning, and quantify improvements in hallucination reduction.
3. Cross-benchmark validation	Test the refined approach on open-source datasets using distributed evaluation to ensure generality.
4. Reward-model adjustment & ablation	Adjust PRM and ORM weightings within GRPO; analyze their effects on factual accuracy and reasoning coherence.
5. Further ablation & exploratory analysis	Perform statistical testing, case analysis, and visualize hallucination emergence during chain-of-thought reasoning.

Reference:

Grounded Hallucination Rates for Top 25 LLMs



Last updated on April 29th, 2025

JOURNAL OF LATEX CLASS FILES, JANUARY 2025

https://arxiv.org/abs/2502.17419

From System 1 to System 2: A Survey of Reasoning Large Language Models

Zhong-Zhi Li*, Duzhen Zhang*, Ming-Liang Zhang§, Jiaxin Zhang§, Zengyan Liu§, Yuxuan Yao§,
Haotian Xu, Junhao Zheng, Pei-Jie Wang, Xiuyi Chen, Yingying Zhang, Fei Yin, Jiahua Dong,
Zhiwei Li, Bao-long Bi, Ling-rui Mei, Jun-Feng Fang, Xiao Liang
Zhijiang Guo†, Le Song†, Cheng-Lin Liu†, Fellow, IEEE

Abstract—Achieving human-level intelligence requires refining the transition from the fast, intuitive System 1 to the slower, more deliberate System 2 reasoning. While System 1 excels in quick, heuristic decisions, System 2 relies on logical reasoning for more accurate judgments and reduced biases. Foundational Large Language Models (LLMs) excel at fast decision-making but lack the depth for complex reasoning, as they have not yet fully embraced the step-by-step analysis characteristic of true System 2 thinking. Recently, reasoning LLMs like OpenAI's o1/o3 and DeepSeek's R1 have demonstrated expert-level performance in fields such as mathematics and coding, closely mimicking the deliberate reasoning of System 2 and showcasing human-like cognitive abilities. This survey begins with a brief overview of the progress in foundational LLMs and the early development of System 2 technologies, exploring how their combination has paved the way for reasoning LLMs. Next, we discuss how to construct reasoning LLMs, analyzing their features, the core methods enabling advanced reasoning, and the evolution of various reasoning LLMs. Additionally, we provide an overview of reasoning benchmarks, offering an in-depth comparison of the performance of representative reasoning LLMs. Finally, we explore promising directions for advancing reasoning LLMs and maintain a real-time GitHub Repository to track the latest developments. We hope this survey will serve as a valuable resource to inspire innovation and drive progress in this rapidly evolving field.

Index Terms—Slow-thinking, Large Language Models, Human-like Reasoning, Decision Making in AI, AGI

1 INTRODUCTION

“Don’t teach. Incentivize.”

—Hyung Won Chung, OpenAI

Version: v1 (major update on February 24, 2025)

*Core contribution, §Significant contribution, †Corresponding author.
Duzhen Zhang, Jiahua Dong, and Le Song are with the Mohamed bin Zayed University of Artificial Intelligence, Abu Dhabi, UAE (E-mail: bladedancer957@gmail.com; dongjiahua1995@gmail.com; le.song@microsoft.com).

Zhong-Zhi Li, Pei-Jie Wang, Bao-Long Bi, Ling-Rui Mei, Xiuyi Chen, Fei Yin, and Zhiwei Li are with the Chinese Academy of Sciences, Beijing, China (E-mail: zhongzhili2023@ia.ac.cn, bilaodong23@ict.ac.cn, melingrui25@ict.ac.cn, xiuyi.chen@gmail.com; fjin@slpr.ac.cn; linc@slpr.ia.ac.cn).

Ming-Liang Zhang is with the Alibaba Group, Beijing, China (E-mail: luxing.zml@taobao.com).
Bao-Long Bi, Ling-Rui Mei are with the Chinese Academy of Sciences, Beijing, China (E-mail: bilaodong23@ict.ac.cn, melingrui25@ict.ac.cn).
Xiao Liang is with the University of California, Los Angeles, USA (E-mail: luxing.zml@ucla.edu).
Zengyan Liu, Yuxuan Yao, Zhiwei Li, and Zhijiang Guo are with City University of Hong Kong and the Hong Kong University of Science and Technology

ACHIEVING human-level intelligence requires refining the transition from *System 1* to *System 2* reasoning [1]–[5]. Dual-system theory suggests that human cognition operates through two modes: *System 1*, which is fast, automatic, and intuitive, enabling quick decisions with minimal effort, and *System 2*, which is slower, more analytical, and deliberate [6], [7]. While *System 1* is efficient for routine tasks, it is prone to cognitive biases, especially in complex or uncertain situations, leading to judgment errors. In contrast, *System 2* relies on logical reasoning and systematic thinking, resulting in more accurate and rational decisions [8]–[11]. By mitigating the biases of *System 1*, *System 2* provides a more refined approach to problem-solving [12]–[15].

The development of foundational Large Language Models (LLMs)¹ has marked a major milestone in Artificial Intelligence (AI). Models such as GPT-4o [16] and DeepSeek-v3 [17] have demonstrated impressive capabilities in text generation, language translation, and a variety of perception tasks [18]–[28]. These models, trained on extensive datasets and utilizing advanced algorithms, excel in understanding and generating human-like responses. However, despite

JOURNAL OF LATEX CLASS FILES, JANUARY 2025

TABLE 6
Statistics of benchmarks for reasoning LLMs.

Domain	Benchmark	Question Type	Venue	Language	Size	Level
Math	AlME 2024 [30]	Open-End	Finding 2023	English	30	Competition
	BBP 2024 [31]	Hybrid	Finding 2023	English	23	Challenging
	MATH-500 [32]	Open-End	Finding 2023	English	500	Competition
	AMC 2023 [33]	Open-End	ACL 2024	English/Chinese	30	Competition
	OpenMath 2023 [34]	Open-End	ACL 2024	English	236	Competition
	Putnam-AXIOM [32]	Open-End	NeurIPS 2024	English	236	Competition
	PRIMOBK [35]	Open-End	ICLR 2024	English	800,000	Hybrid
	Primo 2024 [36]	Open-End	ArXiv 2024	English	1024	Expert
	ProcessBench [31]	Open-End	ArXiv 2024	English	3400	Competition
	LiveBench [31]	Open-End	CLRS 2025	English	-	Frequently Updated
	AI4Maths [37]	Open-End	Humanizing Race	English	13	Competition
	ThinkBench [36]	Hybrid	ArXiv 2025	English	2,912	Expert
	MATH-Portfolios [37]	Open-End	ArXiv 2025	English	276	Competition
	ZetaMath [38]	Open-End	ArXiv 2025	English	1,000	Hybrid
	QuetionSet [39]	Choice	ArXiv 2025	English	38,882	Hybrid
	Math Roll [30]	Open-End	ArXiv 2025	English	-	High School
	CrossBench [32]	Open-End	ArXiv 2025	English	-	Middle School
Code	Codeforces [40]	Open-End	Science 2022	English	-	Expert
	CodeContent [32]	Open-End	ICLR 2024	English	13,610	Competition
	SWIN 2024 [32]	Open-End	ICLR 2024	English	2,994	Expert
	LiveCodeBench [34]	Open-End	ArXiv 2024	English	-	Expert
	CodeCodeBench [32]	Hybrid	ArXiv 2024	English	-	Expert
Science	Coq’s Diamond [36]	Open-End	CLRS 2024	English	448	University
	ML4Bench [37]	Hybrid	NeurIPS 2024	English	5,975	Hybrid
	MM4L 2024 [38]	Choice	NeurIPS 2024	English	12,032	Hybrid
	ML4Bench [37]	Open-End	ArXiv 2024	English	210	Expert
	RewardBench [30]	Open-End	ArXiv 2024	English	-	Hybrid
	QCoDe 2024 [39]	Open-End	QCoDe 2024	English	5,975	Hybrid
	ML4Bench [37]	Open-End	NeurIPS 2024	English	5,975	Hybrid
	JudgeGPT [33]	Open-End	ACL 2024	English	57	University
	THP 2024 [38]	Open-End	ArXIV 2025	English/Chinese	796	Competition
	Prob4Bench [36]	Open-End	ArXIV 2025	English	2,400	Hybrid
	EquationBench [36]	Open-End	ArXIV 2025	English	26,529	University
	SuperGPTQA [38]	Open-End	ArXIV 2025	English	-	Expert
	Sys4Bench [39]	Open-End	ArXIV 2025	English	6,216	Hybrid
	PoP4Bench [40]	Open-End	ArXIV 2025	English	-	Expert
	Doll4Bench [31]	Open-End	ArXIV 2025	English	-	Expert
	FIDEREASON [32]	Open-End	ArXIV 2025	English	-	Expert
Agent	CoCoBench [41]	Open-End	ArXIV 2025	Symbolic	1,000	Expert
	Web4Bench [34]	Open-End	NeurIPS 2022	English	1,600	Hybrid
	SciWorld [34]	Open-End	EMNLP 2022	English	7,200	Hybrid
	Web4Bench [34]	Open-End	CLRS 2024	English	812	Hybrid
	TextCraft [42]	Open-End	NAACL 2023	English	290	Hybrid
	Overworld [38]	Open-End	NeurIPS 2024	English	369	Hybrid
	CoCoBench [41]	Open-End	ArXIV 2025	English	-	Hybrid
	ML4Bench [37]	Open-End	ArXIV 2025	English	-	Competition
	Tox4Comp [31]	Open-End	ArXIV 2025	English	-	Hybrid
	Mobile4Bench [39]	Open-End	ArXIV 2025	English	25	Hybrid
	Text4World [33]	Open-End	ArXIV 2025	English	-	Hybrid
	Web4Games [34]	Open-End	ArXIV 2025	English	50	Hybrid
	UV4Bench [34]	Open-End	ArXIV 2025	English	128	Hybrid
Medicine	JAMA Clinical [36]	Choice	NAACL 2025	English	1,524	Expert
	MedBench [36]	Choice	NAACL 2025	English	6,007	Expert
	Med4Bench [38]	Open-End	ArXIV 2024	English/Chinese	3,848	Expert
	MedOpener [39]	Choice	ArXIV 2025	English	4,460	Expert